

Auralynq Offline Benchmark Report (Sample)

This is a sample report included in Auralynq's public demo corpus, formatted to mirror the structure that **make eval** and **make bench** write to **reports/**. The values below were produced by an earlier run of those exact commands in the fully-offline \$0 configuration (hash embeddings, in-memory vector store, extractive LLM) over a frozen 5-item golden set, and are also published in the project README's Benchmarks section. They illustrate report layout and are not a live, current benchmark of any particular deployment.

Retrieval comparison (k=6, nDCG@10)

Metric	naive	hybrid	PathRAG	full agentic
Recall@k	1.00	1.00	0.80	0.80
nDCG@10	0.900	0.886	0.800	0.800
MRR	0.867	0.850	0.800	0.800
Precision@k	0.167	0.167	0.133	0.133
Latency p50 (ms)	0.1	1.3	0.1	16.6

Qdrant quantization trade-off (289 vectors, dim 256)

Quantization	Recall@10	Memory	Compression
none (fp32)	1.00	289 KB	1x
scalar (int8)	1.00	72 KB	4x
binary (1-bit)	0.50	9 KB	32x

Source: reports/ (make bench output), also published in README.md - Benchmarks. Regenerate current numbers with: make eval && make bench.